Joe Doe

Senior Data Scientist

CONTACT

- joe.doe@gmail.com
- +1 (415) 829-4371
- San Francisco, CA
- linkedin.com/in/joedoe-ds
- github.com/joedoe-ml
- joedoe.dev

SKILLS

Machine Learning	Intermediate
Programming	Intermediate
Data Engineering	Intermediate
MLOps & Cloud	Intermediate
Analytics	Intermediate
SQL (advanced)	Proficient
Tableau/Looker	Proficient
R	Proficient
Statistical Modeling	Advanced
A/B Experimentation	Advanced
Data Mining	Advanced
Product Analytics	Advanced

LANGUAGES

English
Native

SUMMARY

Data Scientist with 6+ years of experience applying quantitative analysis, experimentation, and statistical modeling to deliver product impact at scale for products serving billions of users. Proven track record influencing product strategy and investment decisions through data-driven storytelling and rigorous analytical frameworks, collaborating cross-functionally across Product, Engineering, and Research teams within fast-moving technology environments. Led initiatives leveraging ML and deep learning that generated $40M+ in incremental revenue, with published research at NeurIPS and ICML. Skilled in SQL, Python, and data mining to identify opportunities, measure product success, forecast key metrics, and drive roadmaps aligned with Meta's mission of real-world AI impact at massive scale.

EXPERIENCE

Senior Data Scientist
Mar 2022 — Present

Google · Mountain View, US

Search & AI division — ranking and personalization team

- Designed and deployed a transformer-based query understanding model that improved search relevance by 12%, impacting 2B+ daily queries
- Built an end-to-end MLOps pipeline using Vertex AI, reducing model deployment time from 2 weeks to 4 hours
- Led a team of 4 data scientists to develop a real-time user intent prediction system, increasing ad click-through rate by 8.5%
- Collaborated with product and engineering teams to A/B test 15+ ML features, driving $28M in annual incremental revenue

Data Scientist
Jan 2020 — Feb 2022

Spotify · New York, US

Personalization & recommendations team

- Developed a collaborative filtering model for podcast recommendations, increasing user engagement by 23% across 400M+ users
- Implemented a real-time feature engineering pipeline using Apache Kafka and Spark, processing 50M+ events per hour
- Created an NLP-based content moderation system that reduced manual review workload by 60%
- Mentored 3 junior data scientists and established best practices for experiment design and causal inference

Machine Learning Engineer
Jun 2018 — Dec 2019

DataRobot · Boston, US

AutoML platform — model optimization team

French
Professional Working

Mandarin
Elementary

- Built automated feature selection algorithms that improved model accuracy by 15% on average across 200+ customer datasets
- Developed a time-series forecasting module supporting ARIMA, Prophet, and LSTM architectures
- Reduced model training time by 40% through distributed computing optimizations with Dask and Ray

EDUCATION

Massachusetts Institute of Technology (MIT)

Sep 2016 — May 2018

Master of Science • Computer Science — Machine Learning

University of California, Berkeley

Sep 2012 — May 2016

Bachelor of Science • Statistics & Data Science

PROJECTS

Real-Time Fraud Detection Engine

Jan 2023 — Jun 2023

Built a gradient-boosted ensemble model processing 10K+ transactions per second with 99.7% accuracy and <50ms latency, preventing $12M in annual fraud losses.

LLM-Powered Document Intelligence

Aug 2023 — Jan 2024

Developed a RAG pipeline using GPT-4 and vector embeddings to automate contract analysis, reducing legal review time by 75% for enterprise clients.

Customer Churn Prediction Platform

May 2021 — Nov 2021

Designed an end-to-end ML platform predicting subscriber churn with 91% recall, enabling targeted retention campaigns that reduced churn by 18%.

CERTIFICATIONS

AWS Certified Machine Learning — Specialty

Amazon Web Services

Apr 2023

Google Professional Machine Learning Engineer

Google Cloud

Nov 2022

Deep Learning Specialization

Coursera — Andrew Ng

Mar 2020

PUBLICATIONS

Efficient Attention Mechanisms for Long-Document Understanding

NeurIPS 2023

Dec 2023

Proposed a sparse attention variant achieving 95% of full attention performance with 60% less compute.

Causal Inference in Large-Scale Recommendation Systems
ICML 2022
Jul 2022
Introduced a debiasing framework for recommendation models, improving fairness metrics by 30%.